

1. General coding instructions (*non-negotiable*)

1.1. SQL + DAL discipline

- **SQL queries are first-class citizens:** every DB operation must be a named query constant in the relevant QRY_* class.
- WorkFlowDAL is the only DB entry point for the engine.
 - ✓ Engine/components **must not** call leaf DALs directly (read/write DALs are internal implementation details behind IWorkFlowDAL).
- Adding new DB capability:
 1. Add query constant to the relevant QRY_* class
 2. Extend the relevant DAL(s)
 3. Expose via IWorkFlowDAL

1.2. Formatting rules

- **Models:** keep constructors minimal; prefer object initializer style:
 - ✓ **Preferred:** new X { A = a, B = b }
 - ✓ **Not Preferred:** “everything in constructor”
- Method signatures and constructors must be in one line (no multi-line wrapping).
- SQL query strings must be in one line.

1.3. Transactions + cancellation (DbExecutionLoad)

- DbExecutionLoad is just: (CancellationToken + TransactionHandler).
- Public engine APIs accept **CancellationToken**.
- Internal DB methods accept **DbExecutionLoad**.
- Transaction pattern (engine-owned):

```
var transaction = _dal.CreateNewTransaction();

using var tx = transaction.Begin(false);

var load = new DbExecutionLoad(ct, transaction);
```

- ✓ Commit/Rollback guarded by a committed flag.

1.4. “Ensure/Exists then Insert” rule

- For “upsert” style operations, prefer:
 - ✓ EXISTS → return id
 - ✓ else INSERT → return id
- **Avoid the pattern:** `id = LAST_INSERT_ID(id)` / mutating primary key semantics.

1.5. Naming conventions (engine project)

- Flag enums are prefixed: LifeCycleStateFlag, LifeCycleTransitionFlag, LifeCycleInstanceFlag, LifeCycleTransitionLogFlag, etc.
- DAL group classes are internal sealed (except the few agreed public ones like WorkflowDAL / DALUtil).
- Query placeholders when performing join queries: use the newer placeholders; Example: Joining two tables, where both table has ‘name’ field. In such case, use unique identifier for the field.
 - ✓ ENV_NAME, DEF_NAME (instead of reusing DISPLAY_NAME/NAME where they reduce readability)

2. Concept (high level)

Objective (what this engine is)

A dumb, reliable, DB-backed workflow/state engine that:

- Transitions an instance through states based on triggers (macro workflow). Triggers are nothing by event-contracts raised from the application. Applications doesn't know anything about transition or states. It has awareness only about the events.
- Emits outbound events for:
 - ✓ lifecycle transitions (macro)
 - Application has to decide which internal method to raise when a lifecycle transition is raised.
 - App should also take responsibility of sending the acknowledgements back to the engine.
 - Events to internal methods are **hard-coded** inside the application.
 - ✓ hooks/work-items (micro orchestration)
 - Directly tells the application which internal method to raise.
 - Benefit is that the orchestration can be controlled via json configs.
 - **No hard-coding** of events to internal methods.
- Tracks delivery/processing via ACKs and re-sends until completed
- Automatically triggers timeout events for stale instances
- Produces timeline/reporting; runtime tables are mainly for app activity/status storage

Non-goal: business logic. Business meaning lives in the application + JSON policies/routes, not inside the engine.

3. Core runtime flow (*application → engine → application*)

3.1. Trigger (application → engine)

Application calls:

- `TriggerAsync(LifeCycleTriggerRequest req, ct)`

Minimum req fields:

- `env_code, def_name`; *to find the target definition version*
- `external_ref` (application correlation key); *important as this is how both app and engine correlates*
- `event` (name or code); *to initiate the transition*
- `request_id` (recommended for idempotency);
- optional: `actor, payload`; *actor is for timeline tracking & payload is for carryforward to next stage*
- optional: `ack required` (if you keep this as a request flag); *default true, so we get ack*

3.2. Engine evaluates + persists (engine → DB)

Engine does, inside `one transaction`:

1. Load blueprint (`IBlueprintManager`) — cached read.
2. Ensure instance exists (`IStateMachine.EnsureInstanceAsync`)
3. Apply transition (`IStateMachine.ApplyTransitionAsync`)
 - a. validates transition exists for (`from_state, event`)
 - b. updates instance current state using CAS (*Compare And Set; basically Upsert*)
 - c. inserts lifecycle timeline (`lifecycle, lifecycle_data`)
4. Resolve policy (`IPolicyEnforcer.ResolvePolicyAsync`)
 - a. optionally attach policy to instance
5. Emit hooks (`IPolicyEnforcer.EmitHooksAsync`)
 - a. inserts hook rows (work items) Note: “*Emit hooks*” = *derive hook work-items from policy/routes and persist them, so they can be dispatched reliably (with ACK tracking), same as lifecycle events.*
 - b. Prepare Hook Emissions in memory. (Emit hooks doesn’t raise events immediately. Not inside the transaction.)

6. Create ACK records (**IAckManager**)
 - a. lifecycle ack (linked to lifecycle_id)
 - b. hook ack (linked to hook_id)
 - c. consumer tracking rows per consumer

Commit first.

Only after commit, the engine raises C# events.

3.3. Emit outbound events (engine → application)

The engine exposes **two C# events** only:

1. **EventRaised**(**ILifeCycleEvent**) — actionable
 - a. used for both lifecycle transition events and hook events
2. **NoticeRaised**(**LifeCycleNotice**) — informational
 - a. retries, failures, duplicates, stale warnings, etc.

Base event requirement (mandatory):

All actionable events (transition + hook) derive from one base event and MUST include:

- external_ref
- ack_guid
- plus event information (transition/hook details)

3.4. Ack updates (application → engine)

Application calls:

- **AckAsync**(consumerId, ackGuid, outcome, ...)

ACK stages: (**AckOutcome**)

- **Delivered**: “I received the event” (fast); *expected in less than 30 seconds*.
- **Processed**: “I finished processing the event” (slow); *expected in less than 5 minutes*.
- Failed/Retry can exist as outcomes, but core persistence today is status + retry_count/last_retry.

AckOutcome vs **AckStatus**

- **AckOutcome** is what the application reports to the Engine. *i.o.w; Input from Application*
- **AckStatus** is what the database stores as a state of an ack record. *i.o.w; Persisted state*

Mapping:

- Outcome.Delivered → Status.Delivered
- Outcome.Processed → Status.Processed
- Outcome.Failed → Status.Failed
- Outcome.Retry → Status.Pending + MarkRetry(...) (increment retry_count/last_retry)

3.5. Reliability: Monitor (the heart of idempotency)

Monitor runs periodically and does **two major jobs**:

A) ACK re-dispatch (reliable delivery)

Continuously queries ack dispatch views and re-raises events:

1. Pending too long (delivery not confirmed)

- If an ack is Pending for more than **30 seconds**:
 - ✓ re-raise the same event again (same ack_guid, same payload)
 - ✓ mark retry (retry_count, last_retry)
 - ✓ raise a notice containing ackGuid + instance id + consumer id + kind

2. Delivered but not processed too long

- If an ack is Delivered but not Processed for more than **5 minutes**:
 - ✓ re-raise/remind (same event, same ack_guid)
 - ✓ mark retry
 - ✓ raise notice (ackGuid + instance id + consumer id + kind)

These timings are captured in `WorkFlowEngineOptions`.

B) Stale state automation (timeout events)

Monitor also scans instances:

- If an instance remains in a state beyond **its configured timeout_minutes**:
 - ✓ raise notice: stale instance detected
 - ✓ if blueprint defines timeout_event for that state:
 - auto-trigger that event as system-triggered (same pipeline as normal trigger)

4. Examples (end-to-end)

4.1. Example A — Normal transition with ACK

1. App triggers

- ✓ Env: 1, Def: "VendorPreQualification"
- ✓ ExternalRef: "VENDOR-00042"
- ✓ Event: "Submit"
- ✓ RequestId: "req-2026-01-04-0001"
- ✓ Payload: { ... }

2. Engine commits DB

- ✓ Ensures instance (def_version_id, external_ref)
- ✓ Current state moves: Draft → Submitted
- ✓ Inserts lifecycle row + lifecycle_data
- ✓ Creates one lifecycle ack and consumer rows (for resolved consumers)

3. Engine raises transition event

Application receives ILifeCycleEvent with:

- ✓ external_ref = "VENDOR-00042"
- ✓ ack_guid = "..." critical
- ✓ transition details: from_state, to_state, event, lifecycle_id, occurred_at
- ✓ optional policy info

4. App acks Delivered

- ✓ AckAsync(consumerId, ackGuid, Delivered)

5. App finishes work and acks Processed

- ✓ AckAsync(consumerId, ackGuid, Processed)

4.2. Example B — Pending ACK retry (30 seconds)

1. Engine dispatched event at T0 but app didn't ack Delivered.
2. At T0 + 30s, monitor finds:

- ✓ status = Pending
 - ✓ last_retry older than threshold
3. Monitor re-raises the same event (same ack_guid)
 4. Monitor raises notice:
 - ✓ Kind = AckRetryPending
 - ✓ includes: ack_guid, instance_id, external_ref, consumer_id, retry_count

4.3. Example C — Delivered but not processed reminder (5 minutes)

1. App acks Delivered quickly, but processing takes too long.
2. At Delivered + 5m, monitor re-raises event as a reminder.
3. Notice raised:
 - ✓ Kind = AckReminderProcessedPending
 - ✓ includes same fields + age

4.4. Example D — Stale state timeout automation

1. Instance enters state Review with:
 - ✓ timeout_minutes = 60
 - ✓ timeout_event = "AutoReject"
2. Monitor detects instance still in Review after > 60 minutes.
3. Monitor raises notice:
 - ✓ Kind = StateStale
 - ✓ includes: instance_id, external_ref, state, age, timeout_event
4. Monitor triggers system event "AutoReject" through the same transactional trigger pipeline.

5. Major engine components and roles

5.1. WorkFlowDAL (DB gateway)

Role: Single entry point to DB for all engine operations.

- Consolidates all read/write DALs
- Enforces query-first discipline (QRY_*)
- Exposes CreateNewTransaction() for engine-owned transactions

5.2. BlueprintImporter (IBlueprintImporter)

Role: Import-time canonicalization (JSON → SQL).

- Loads definition JSON: states/events/transitions/categories/timeouts
- Loads policy JSON: routes/hooks/emit rules
- Writes normalized SQL rows:
 - ✓ environment, definition, def_version
 - ✓ state, events, transition
 - ✓ policy, def_policy link (and any route tables you maintain)
- Should run entire import in one transaction.

5.3. BlueprintManager (IBlueprintManager)

Role: Runtime blueprint read + cache.

- Builds LifeCycleBlueprint from DB (states/events/transitions lookup maps)
- Caches latest by (env_code, def_name) and by def_version_id
- Supports invalidation after imports
- Important cache rule: do not “poison” cache with a canceled task (avoid capturing ct inside cached lazy).

5.4. StateMachine (IStateMachine)

Role: Pure state transition executor (dumb by design).

- Ensures instance exists for (def_version_id, external_ref)

- Resolves event (by code or name)
- Validates transition exists
- Updates instance state using CAS to avoid race conflicts
- Writes lifecycle log/timeline (lifecycle, lifecycle_data)
- Returns ApplyTransitionResult (Applied/Reason/From/To/Event/LifeCycleId)

5.5. PolicyEnforcer (IPolicyEnforcer)

Role: Routes/policy evaluation + hook emission.

- Resolves policy for a given state after transition
- Emits hooks/work-items when policy says so:
 - ✓ creates hook rows in DB
 - ✓ prepares hook emission objects with metadata/payload
- Does not implement business logic; only reads policy JSON and translates to DB + emitted events.

5.6. AckManager (IAckManager)

Role: ACK persistence + consumer tracking + monitor-ready dispatch lists.

- Creates ack records and links:
 - ✓ lc_ack (lifecycle_id → ack_id)
 - ✓ hook_ack (hook_id → ack_id)
- Creates consumer status rows (ack_consumer)
- Updates statuses on AckAsync calls from the app
- Provides monitor methods to list “pending dispatch” items from ack dispatch queries

Critical rule: ACK identifiers are stable:

- re-dispatch always uses the same ack_guid for the same lifecycle/hook item.

5.7. LifeCycleMonitor (ILifeCycleMonitor) — scheduler

Role: Timer + safety gate.

- Runs on a periodic interval (PeriodicTimer)
- Guarantees single active execution (Interlocked gate)

- Calls `WorkFlowEngine.RunMonitorOnceAsync(ct)` each tick; (*Very crucial concept. Monitor itself is generic and doesn't have any business logic associated. The logic lies with in the workflow engine itself*)
- Handles monitor exceptions (typically by raising a notice)

5.8. WorkFlowEngine (IWorkFlowEngine, also implements ILifeCycleEngine)

Role: The orchestrator + API surface.

- Owns all core components:
 - ✓ BlueprintManager, Importer, StateMachine, PolicyEnforcer, AckManager, RuntimeEngine, Monitor
- Exposes the public operations:
 - ✓ TriggerAsync, AckAsync, cache invalidation, and monitor run
- Raises only two C# events:
 - ✓ actionable events (EventRaised)
 - ✓ informational notices (NoticeRaised)
- Owns transactions for Trigger:
 - ✓ **DB changes in transaction**
 - ✓ **commit**
 - ✓ **then raise outbound events** (failures become notices; monitor re-sends)

5.9. RuntimeEngine (IRuntimeEngine)

Role: Observability + app activity/status storage (**not “smart workflow”**).

- Allows application to store activity/runtime rows:
 - ✓ statuses, frozen, lc_id linkage, payload/data
- Supports timeline/reporting
- Does not participate in workflow correctness except for reporting.